

Lecture 4:

ELF, Compiler, Linker, and Loader

Sanchuan Chen

schen@auburn.edu

8/31/2023



Executable and Linkable Format

- ▶ Standard binary format for object files
- ▶ Derived from AT&T System V Unix, now supported by UNIX/Linux
- ▶ A unified format for
 - ▶ **Relocatable object files (.o)**: Created by compilers or assemblers. Need to be processed by the linker before running.
 - ▶ **Executable object files**: Have all relocation done and all symbol resolved except perhaps shared library symbols that must be resolved at run time.
 - ▶ **Shared object files (.so)**: Shared library containing both symbol information for the linker and directly runnable code for run time

More in the specification: <https://refspecs.linuxfoundation.org/elf/elf.pdf>

Example (hexdump -C a.out && readelf -e a.out)

00000000	7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00	
00000010	02 00 03 00 01 00 00 00 80 9a 04 08 34 00 00 00	ELF Header
00000020	78 64 01 00 00 00 00 00 34 00 20 00 08 00 28 00	
00000030	1c 00 1b 00 06 00 00 00 34 00 00 00 34 80 04 08	
00000040	34 80 04 08 00 01 00 00 00 01 00 00 05 00 00 00	
00000050	04 00 00 00 03 00 00 00 34 01 00 00 34 81 04 08	Program Header
00000060	34 81 04 08 13 00 00 00 13 00 00 00 04 00 00 00	
	...	
00000130	04 00 00 00 2f 6c 69 62 2f 6c 64 2d 6c 69 6e 75	sec[1].interp
00000140	78 2e 73 6f 2e 32 00 00 04 00 00 00 10 00 00 00	sec[2].note.ABI-tag
00000150	01 00 00 00 47 4e 55 00 00 00 00 02 00 00 00 00	
00000160	06 00 00 00 08 00 00 00 61 00 00 00 68 00 00 00	...
	...	
000014b0	ff 35 08 e1 05 08 ff 25 0c e1 05 08 00 00 00 00	
000014c0	ff 25 10 e1 05 08 68 00 00 00 00 e9 ff ff ff	sec[12].plt
	...	
00001a70	ff 25 7c e2 05 08 68 d8 02 00 00 e9 30 fa ff ff	sec[13].text
00001a80	31 ed 5e 89 e1 83 e4 f0 50 54 52 68 60 9e 05 08	...
	...	
00016390	6c 73 00 00 52 a8 bd 5a 00 2e 73 68 73 74 72 74	sec[27].shstrtab
000163a0	61 62 00 2e 69 6e 74 65 72 70 00 2e 6e 6f 74 65	
	...	
00016470	75 67 6c 69 6e 6b 00 00 00 00 00 00 00 00 00 00	
00016480	00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00	
	...	
000164a0	0b 00 00 00 01 00 00 00 02 00 00 00 34 81 04 08	
000164b0	34 01 00 00 13 00 00 00 00 00 00 00 00 00 00 00	Section Header
	...	
000168c0	98 63 01 00 df 00 00 00 00 00 00 00 00 00 00 00	
000168d0	01 00 00 00 00 00 00 00	

Section Headers:										
[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
[0]		NULL	00000000	000000	000000	00		0	0	0
[1]	.interp	PROGBITS	08048134	000134	000013	00	A	0	0	1
[2]	.note.ABI-tag	NOTE	08048148	000148	000020	00	A	0	0	4
		...								
[12]	.plt	PROGBITS	080494b0	0014b0	0005d0	04	AX	0	0	4
[13]	.text	PROGBITS	08049a80	001a80	0104bc	00	AX	0	0	16
[14]	.fini	PROGBITS	08059f3c	011f3c	00001c	00	AX	0	0	4
[15]	.rodata	PROGBITS	08059f60	011f60	003e4c	00	A	0	0	32
		...								
[27]	.shstrtab	STRTAB	00000000	016398	0000df	00		0	0	1

ELF header

- Magic number, type (.o, exec, .so), machine, byte ordering, etc.

```
typedef struct elf32_hdr{
    unsigned char e_ident[EI_NIDENT];
    Elf32_Half e_type; /* ET_EXEC ET_DYN */
    Elf32_Half e_machine;
    Elf32_Word e_version;
    Elf32_Addr e_entry; /* Entry point */
    Elf32_Off e_phoff;
    Elf32_Off e_shoff;
    Elf32_Word e_flags;
    Elf32_Half e_ehsize;
    Elf32_Half e_phentsize;
    Elf32_Half e_phnum;
    Elf32_Half e_shentsize;
    Elf32_Half e_shnum;
    Elf32_Half e_shstrndx;
} Elf32_Ehdr;
```

.interp

- ▶ An ELF interpreter, normally `/lib/ld-linux.so.2`

Program header table

- ▶ page size, virtual addresses for memory segments (sections), segment sizes.

.text section

- ▶ code

.data section

- ▶ initialized (static) data

.bss section

- ▶ uninitialized (static) data
- ▶ “Block Started by Symbol” / “Better Save Space”
- ▶ has section header but occupies no space

.symtab section

- ▶ symbol table
- ▶ procedure and static variable names
- ▶ section names

.rel.text section


- ▶ relocation info for .text section
- ▶ addresses of instructions that will need to be modified in the executable
- ▶ instructions for modifying.

.rel.data section

- ▶ relocation info for .data section
- ▶ addresses of pointer data that will need to be modified in the merged executable

.debug section

- ▶ info for symbolic debugging (gcc -g)

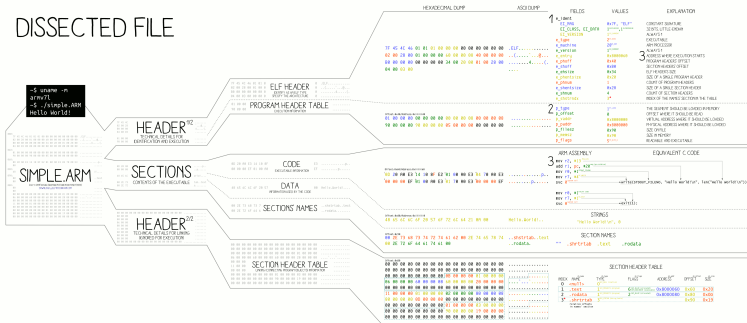


- ▶ A single **segment** usually consist of **several sections**.
 - ▶ E.g., a loadable read-only segment could contain sections for executable code, read-only data, and symbols for the dynamic linker.
- ▶ Sections are intended for further processing by the **linker**, while the segments are intended to be mapped into memory by the **loader**.

ELF¹⁰¹ a Linux executable walk-through

ANGE ALBERTINI
CORKAM.COM

DISSECTED FILE



LOADING PROCESS

1. HEADER


THE ELF HEADER IS PARSED
THE PROGRAM HEADER IS PARSED
(SECTIONS ARE NOT USED)

2. MAPPING

THE FILE IS MAPPED IN MEMORY
ACCORDING TO ITS SEGMENT(S)

3. EXECUTION

ENTRY IS CALLED
SYSCALLSTM ARE ACCESSED VIA:
- SYSCALL NUMBER IN THE R7 REGISTER
- CALLING INSTRUCTION SVC



TRIVIA

THE ELF WAS FIRST SPECIFIED BY U.S. "C" AND "U"™
FOR UNIX SYSTEM V, IN 1989

THE ELF IS USED, AMONG OTHERS, IN:
- LINUX, ANDROID, BSD, SOLARIS, BEOS
- PSP, PLAYSTATION 2-4, DREAMCAST, GAMECUBE, Wii
- VARIOUS OSes MADE BY SAMSUNG, ERICSSON, NOKIA,
- MICROCONTROLLERS FROM ATMEL, TEXAS INSTRUMENTS

Why Compiler?

- ▶ Compiler parses the source code
- ▶ Compiler knows exactly the behavior of the program
- ▶ Compiler has to know the machine details (when generating the code)
- ▶ Compiler-based security solutions, e.g., **vulnerability detection**
 - ▶ Stack overflow
 - ▶ Integer overflow
 - ▶ Heap overflow
 - ▶ format string
 - ▶ Double free
 - ▶ Use-after-free

GNU Compiler Collection (GCC)

- ▶ **Developer(s)** GNU Project
- ▶ **Initial release** May 23, 1987
- ▶ **Newest version** GCC 13.2
- ▶ **Written** in C, C++
- ▶ **Operating system** Cross-platform
- ▶ **Platform** GNU
- ▶ **Type** Compiler
- ▶ **License** GNU General Public License (version 3 or later)
- ▶ **Website** gcc.gnu.org

Compilation system includes the phases

- ▶ Preprocessor
- ▶ Compiler
 - ▶ e.g., Optimizer
- ▶ Assembler
- ▶ Linker

Compiler Driver coordinates these phases

GCC(1)

NAME

gcc - GNU project C and C++ compiler

SYNOPSIS

```
gcc [-c|-S|-E] [-std=standard]
    [-g] [-pg] [-Olevel]
    [-Wwarn...] [-Wpedantic]
    [-Idir...] [-Ldir...]
    [-Dmacro[=defn]...] [-Umacro]
    [-foption...] [-mmachine-option...]
    [-o outfile] [@file] infile...
```

DESCRIPTION

When you invoke GCC, it normally does preprocessing, compilation, assembly and linking.

Compiler Driver command example

```
schen@linux:~/comp6700/lec04$ gcc -m32 -fno-pic helloworld.c -v
Using built-in specs.
COLLECT_GCC=gcc
COLLECT_LTO_WRAPPER=/usr/lib/gcc/x86_64-linux-gnu/9/lto-wrapper
OFFLOAD_TARGET_NAMES=nvptx-none:hsa
OFFLOAD_TARGET_DEFAULT=1
Target: x86_64-linux-gnu
Configured with: ../src/configure -v --with-pkgversion='Ubuntu 9.3.0-17ubuntu1~20.04'
--with-bugurl=file:///usr/share/doc/gcc-9/README.Bugs
--enable-languages=c,ada,c++,go,brig,d,fortran,objc,obj-c++,gm2 --prefix=/usr
--with-gcc-major-version-only --program-suffix=-9 --program-prefix=x86_64-linux-gnu-
--enable-shared --enable-linker-build-id --libexecdir=/usr/lib
--without-included-gettext --enable-threads=posix --libdir=/usr/lib --enable-nls
--enable-clocale=gnu --enable-libstdcxx-debug --enable-libstdcxx-time=yes
--with-default-libstdcxx-abi=new --enable-gnu-unique-object --disable-vtable-verify
--enable-plugin --enable-default-pie --with-system-zlib
--with-target-system-zlib=auto --enable-objc-gc=auto --enable-multiarch
--disable-werror --with-arch=32=i686 --with-abi=m64
--with-multilib-list=m32,m64,mx32 --enable-multilib --with-tune=generic
--enable-offload-targets=nvptx-none=/build/gcc-9-HskZEa/gcc-9-9.3.0/debian/tmp-nvptx
/usr,hsa --without-cuda-driver --enable-checking=release --build=x86_64-linux-gnu
--host=x86_64-linux-gnu --target=x86_64-linux-gnu
Thread model: posix
gcc version 9.3.0 (Ubuntu 9.3.0-17ubuntu1~20.04)
COLLECT_GCC_OPTIONS='-m32' '-fno-pic' '-v' '-mtune=generic' '-march=i686'
/usr/lib/gcc/x86_64-linux-gnu/9/cc1 -quiet -v -imultilib 32 -imultiarch
i386-linux-gnu helloworld.c -quiet -dumpbase helloworld.c -m32 -mtune=generic
-march=i686 -auxbase helloworld -version -fno-pic -fasynchronous-unwind-tables
-fstack-protector-strong -Wformat -Wformat-security -fstack-clash-protection
-fcf-protection -o /tmp/ccWCdCj6.s
GNU C17 (Ubuntu 9.3.0-17ubuntu1~20.04) version 9.3.0 (x86_64-linux-gnu)
compiled by GNU C version 9.3.0, GMP version 6.2.0, MPFR version 4.0.2, MPC
version 1.1.0, isl version isl-0.22.1-GMP
```

Compiler Driver command example (continued)

```

GGC heuristics: --param ggc-min-expand=100 --param ggc-min-heapsize=131072
ignoring nonexistent directory "/usr/local/include/i386-linux-gnu"
ignoring nonexistent directory "/usr/lib/gcc/x86_64-linux-gnu/9/include-fixed"
ignoring nonexistent directory
"/usr/lib/gcc/x86_64-linux-gnu/9/../../../../x86_64-linux-gnu/include"
ignoring nonexistent directory "/usr/include/i386-linux-gnu"
#include "..." search starts here:
#include <...> search starts here:
  /usr/lib/gcc/x86_64-linux-gnu/9/include
  /usr/local/include
  /usr/include
End of search list.
GNU C17 (Ubuntu 9.3.0-17ubuntu1~20.04) version 9.3.0 (x86_64-linux-gnu)
  compiled by GNU C version 9.3.0, GMP version 6.2.0, MPFR version 4.0.2, MPC
  version 1.1.0, isl version isl-0.22.1-GMP

GGC heuristics: --param ggc-min-expand=100 --param ggc-min-heapsize=131072
Compiler executable checksum: bbf13931d8de1abe14040c9909cb6969
COLLECT_GCC_OPTIONS='-m32' '-fno-pic' '-v' '-mtune=generic' '-march=i686'
  as -v --32 -o /tmp/ccRHA1r7.o /tmp/ccWCdCj6.s
GNU assembler version 2.34 (x86_64-linux-gnu) using BFD version (GNU Binutils for
Ubuntu) 2.34
COMPILER_PATH=/usr/lib/gcc/x86_64-linux-gnu/9:/usr/lib/gcc/x86_64-linux-gnu/9:/usr
/lib/gcc/x86_64-linux-gnu/9:/usr/lib/gcc/x86_64-linux-gnu/9:/usr/lib/gcc/x86_64-linu
x-gnu/
LIBRARY_PATH=/usr/lib/gcc/x86_64-linux-gnu/9/32:/usr/lib/gcc/x86_64-linux-gnu/9/..
/../../../../i386-linux-gnu:/usr/lib/gcc/x86_64-linux-gnu/9/../../../../lib32:/lib/i386-l
inux-gnu:/lib/../../lib32:/usr/lib/i386-linux-gnu:/usr/lib/../../lib32:/usr/lib/gcc/x8
6_64-linux-gnu/9:/usr/lib/gcc/x86_64-linux-gnu/9/../../../../i386-linux-gnu:/usr/lib/
gcc/x86_64-linux-gnu/9/../../../../lib/i386-linux-gnu:/lib:/usr/lib/i386-linux-gnu/
:/usr/lib/

```

Compiler Driver command example (continued)

```
COLLECT_GCC_OPTIONS='-m32' '-fno-pic' '-v' '-mtune=generic' '-march=i686'
/usr/lib/gcc/x86_64-linux-gnu/9/collect2 -plugin
/usr/lib/gcc/x86_64-linux-gnu/9/liblto_plugin.so
-plugin-opt=/usr/lib/gcc/x86_64-linux-gnu/9/lto-wrapper
-plugin-opt=-fresolution=/tmp/ccQlh2N8.res -plugin-opt=-pass-through=-lgcc
-plugin-opt=-pass-through=-lgcc_s -plugin-opt=-pass-through=-lc
-plugin-opt=-pass-through=-lgcc -plugin-opt=-pass-through=-lgcc_s --build-id
--eh-frame-hdr -m elf_i386 --hash-style=gnu --as-needed -dynamic-linker
/lib/ld-linux.so.2 -pie -z now -z relro
/usr/lib/gcc/x86_64-linux-gnu/9/../../../../lib32/Scrt1.o
/usr/lib/gcc/x86_64-linux-gnu/9/../../../../lib32/crti.o
/usr/lib/gcc/x86_64-linux-gnu/9/32/crtbeginS.o -L/usr/lib/gcc/x86_64-linux-gnu/9/32
-L/usr/lib/gcc/x86_64-linux-gnu/9/../../../../i386-linux-gnu
-L/usr/lib/gcc/x86_64-linux-gnu/9/../../../../lib32 -L/lib/i386-linux-gnu
-L/lib/./lib32 -L/usr/lib/i386-linux-gnu -L/usr/lib/./lib32
-L/usr/lib/gcc/x86_64-linux-gnu/9
-L/usr/lib/gcc/x86_64-linux-gnu/9/../../../../i386-linux-gnu
-L/usr/lib/gcc/x86_64-linux-gnu/9/../../../../lib32 -L/lib/i386-linux-gnu
-L/usr/lib/i386-linux-gnu /tmp/ccRHA1r7.o -lgcc --push-state --as-needed -lgcc_s
--pop-state -lc -lgcc --push-state --as-needed -lgcc_s --pop-state
/usr/lib/gcc/x86_64-linux-gnu/9/32/crtendS.o
/usr/lib/gcc/x86_64-linux-gnu/9/../../../../lib32/crtn.o
COLLECT_GCC_OPTIONS='-m32' '-fno-pic' '-v' '-mtune=generic' '-march=i686'
```


Linker

In computer science, a linker or link editor is a program that takes one or more objects generated by a compiler and combines them into a single executable program. It can happen in various stages.

- ➊ Linking can be done at compile time
- ➋ at load time (by loaders)
- ➌ at run time (by application programs)

Linker

Compilation refers to the processing of source code files (.c, .cc, or .cpp) and the creation of an 'object' file. This step doesn't create anything the user can actually run. Instead, the compiler merely produces the machine language instructions that correspond to the source code file that was compiled.


Linking refers to the creation of a single executable file from multiple object files. In this step, it is common that the linker will complain about undefined functions (commonly, main itself).

Why Linking: (1) Modularity

- ▶ Program can be written as a collection of smaller source files, rather than one monolithic mass.
- ▶ Can build libraries of common functions (more on this later)
 - ▶ e.g., Math library, standard C library

Why Linking: (2) Efficiency

- ▶ Time: Separate compilation
 - ▶ Change one source file, compile, and then relink.
 - ▶ No need to recompile other source files.
- ▶ Space: Libraries
 - ▶ Common functions can be aggregated into a single file...
 - ▶ Yet executable files and running memory images contain only code for the functions they actually use.
 - ▶ Non-shared v.s. Shared Library



```

$ ar -tv libc.a
rw-r--r-- 0/0 2488 Dec 31 19:00 1969 init-first.o
rw-r--r-- 0/0 12992 Dec 31 19:00 1969 libc-start.o
rw-r--r-- 0/0 608 Dec 31 19:00 1969 sysdep.o
rw-r--r-- 0/0 2704 Dec 31 19:00 1969 version.o
rw-r--r-- 0/0 2432 Dec 31 19:00 1969 check_fds.o
rw-r--r-- 0/0 3432 Dec 31 19:00 1969 libc-tls.o
rw-r--r-- 0/0 2184 Dec 31 19:00 1969 elf-init.o
rw-r--r-- 0/0 1144 Dec 31 19:00 1969 dso_handle.o
rw-r--r-- 0/0 1104 Dec 31 19:00 1969 errno.o
rw-r--r-- 0/0 1384 Dec 31 19:00 1969 errno-loc.o
rw-r--r-- 0/0 3152 Dec 31 19:00 1969 iconv_open.o
rw-r--r-- 0/0 3104 Dec 31 19:00 1969 iconv.o
rw-r--r-- 0/0 1480 Dec 31 19:00 1969 iconv_close.o
rw-r--r-- 0/0 4688 Dec 31 19:00 1969 gconv_open.o
rw-r--r-- 0/0 2648 Dec 31 19:00 1969 gconv.o
rw-r--r-- 0/0 1576 Dec 31 19:00 1969 gconv_close.o
rw-r--r-- 0/0 18192 Dec 31 19:00 1969 gconv_db.o
rw-r--r-- 0/0 14408 Dec 31 19:00 1969 gconv_conf.o
...
  
```

main.c

```
extern void swap();

int buf[2] = {1, 2};
int main()
{
    swap();
    return 0;
}
```

swap.c

```
extern int buf[];

int *bufp0 = &buf[0];
static int *bufp1;

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

main.o

```
schen@linux:~/comp6700/lec04$ objdump -r -d main.o
00000000 <main>:
    0: f3 0f 1e fb          endbr32
    4: 55                   push    %ebp
    5: 89 e5               mov     %esp,%ebp
    7: 83 e4 f0           and     $0xffffffff0,%esp
    a: e8 fc ff ff ff     call    b <main+0xb>
    b: R_386_PC32 swap
    f: b8 00 00 00 00     mov     $0x0,%eax
   14: c9                 leave
   15: c3                 ret
```

```
schen@linux:~/comp6700/lec04$ objdump -r main.o
```

```
main.o:      file format elf32-i386
```

```
RELOCATION RECORDS FOR [.text]:
OFFSET      TYPE          VALUE
0000000b R_386_PC32      swap
```

Dynamic Linking

NAME

ld.so, ld-linux.so - dynamic linker/loader

SYNOPSIS

The dynamic linker can be run either indirectly by running some dynamically linked program or shared object (in which case no command-line options to the dynamic linker can be passed and, in the ELF case, the dynamic linker which is stored in the .interp section of the program is executed) ...

test_foo.c

```
#include <stdio.h>

extern int foo(void);

int main () {
    printf("foo returns %d\n",foo());
    return 0;
}
```

foo.c

```
static int num = 6700;

int foo(void) {
    return num;
}
```

test_foo

```
$ gcc -m32 -fPIC -shared -o libfoo.so foo.c
$ gcc -m32 -no-pie -fno-pic -ggdb3 -Wall -Wpedantic
-fno-stack-protector test_foo.c -o test_foo libfoo.so

$ export LD_LIBRARY_PATH=$LD_LIBRARY_PATH$:.

$ ldd test_foo
    linux-gate.so.1 (0xf7ed9000)
    libfoo.so (0xf7ecf000)
    libc.so.6 => /lib/i386-linux-gnu/libc.so.6 (0xf7ccc000)
    /lib/ld-linux.so.2 (0xf7eda000)
```



```

08049030 <.plt>:
08049030: ff 35 04 c0 04 08    pushl 0x804c004
08049036: ff 25 08 c0 04 08    jmp    *0x804c008
0804903c: 0f 1f 40 00          nopl   0x0(%eax)
08049040: f3 0f 1e fb          endbr32
08049044: 68 00 00 00 00      push  $0x0
08049049: e9 e2 ff ff ff      jmp    8049030 <printf@plt-0x40>
0804904e: 66 90               xchg  %ax,%ax
08049050: f3 0f 1e fb          endbr32
08049054: 68 08 00 00 00      push  $0x8
08049059: e9 d2 ff ff ff      jmp    8049030 <printf@plt-0x40>
0804905e: 66 90               xchg  %ax,%ax
08049060: f3 0f 1e fb          endbr32
08049064: 68 10 00 00 00      push  $0x10
08049069: e9 c2 ff ff ff      jmp    8049030 <printf@plt-0x40>
0804906e: 66 90               xchg  %ax,%ax
  
```

Disassembly of section .plt.sec:

```

08049070 <printf@plt>:
08049070: f3 0f 1e fb          endbr32
08049074: ff 25 0c c0 04 08    jmp    *0x804c00c
0804907a: 66 0f 1f 44 00 00    nopw   0x0(%eax,%eax,1)
  
```

```

08049080 <__libc_start_main@plt>:
08049080: f3 0f 1e fb          endbr32
08049084: ff 25 10 c0 04 08    jmp    *0x804c010
0804908a: 66 0f 1f 44 00 00    nopw   0x0(%eax,%eax,1)
  
```

```

08049090 <foo@plt>:
08049090: f3 0f 1e fb          endbr32
08049094: ff 25 14 c0 04 08    jmp    *0x804c014
0804909a: 66 0f 1f 44 00 00    nopw   0x0(%eax,%eax,1)
  
```

```

080491b6 <main>:
080491b6: f3 0f 1e fb          endbr32
080491ba: 8d 4c 24 04          lea    0x4(%esp),%ecx
080491be: 83 e4 f0             and    $0xffffffff0,%esp
080491c1: ff 71 fc             pushl  -0x4(%ecx)
080491c4: 55                  push   %ebp
080491c5: 89 e5               mov    %esp,%ebp
080491c7: 51                  push   %ecx
080491c8: 83 ec 04            sub    $0x4,%esp
080491cb: e8 c0 fe ff ff      call   8049090 <foo@plt>
080491d0: 83 ec 08            sub    $0x8,%esp
080491d3: 50                  push   %eax
080491d4: 68 08 a0 04 08      push   $0x804a008
080491d9: e8 92 fe ff ff      call   8049070 <printf@plt>

```

...

\$ readelf -e test_foo

ELF Header:

```

Magic:    7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00
Class:                                ELF32
Data:                                2's complement, little endian
Version:    1 (current)
OS/ABI:      UNIX - System V
ABI Version:    0
Type:        EXEC (Executable file)
Machine:      Intel 80386
Version:      0x1
Entry point address: 0x80490a0

```

...

Section Headers:

[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
[0]		NULL	00000000	000000	000000	00		0	0	0
[1]	.interp	PROGBITS	080481b4	0001b4	000013	00	A	0	0	1
[2]	.note.gnu.build-id	NOTE	080481c8	0001c8	000024	00	A	0	0	4
[3]	.note.gnu.property	NOTE	080481ec	0001ec	00001c	00	A	0	0	4
[4]	.note.ABI-tag	NOTE	08048208	000208	000020	00	A	0	0	4
[5]	.gnu.hash	GNU_HASH	08048228	000228	000020	04	A	6	0	4
[6]	.dynsym	DYNSYM	08048248	000248	000060	10	A	7	1	4
[7]	.dynstr	STRTAB	080482a8	0002a8	00005a	00	A	0	0	1
[8]	.gnu.version	VERSYM	08048302	000302	00000c	02	A	6	0	2
[9]	.gnu.version_r	VERNEED	08048310	000310	000020	00	A	7	1	4
[10]	.rel.dyn	REL	08048330	000330	000008	08	A	6	0	4
[11]	.rel.plt	REL	08048338	000338	000018	08	AI	6	24	4
[12]	.init	PROGBITS	08049000	001000	000024	00	AX	0	0	4
[13]	.plt	PROGBITS	08049030	001030	000040	04	AX	0	0	16
[14]	.plt.sec	PROGBITS	08049070	001070	000030	10	AX	0	0	16
[15]	.text	PROGBITS	080490a0	0010a0	0001c9	00	AX	0	0	16
[16]	.fini	PROGBITS	0804926c	00126c	000018	00	AX	0	0	4
[17]	.rodata	PROGBITS	0804a000	002000	000018	00	A	0	0	4
[18]	.eh_frame_hdr	PROGBITS	0804a018	002018	00004c	00	A	0	0	4
[19]	.eh_frame	PROGBITS	0804a064	002064	00011c	00	A	0	0	4
[20]	.init_array	INIT_ARRAY	0804bf04	002f04	000004	04	WA	0	0	4
[21]	.fini_array	FINI_ARRAY	0804bf08	002f08	000004	04	WA	0	0	4
[22]	.dynamic	DYNAMIC	0804bf0c	002f0c	0000f0	08	WA	7	0	4
[23]	.got	PROGBITS	0804bffc	002ffc	000004	04	WA	0	0	4
[24]	.got.plt	PROGBITS	0804c000	003000	000018	04	WA	0	0	4
[25]	.data	PROGBITS	0804c018	003018	000008	00	WA	0	0	4
[26]	.bss	NOBITS	0804c020	003020	000004	00	WA	0	0	1

Contents of section .got.plt:

```
804c000 0cbf0408 00000000 00000000 40900408 .....@...
804c010 50900408 60900408 .....P...'
```

Loader

Compiler: A language translator that converts a complete program into machine language to produce a program that the computer can process in its entirety.

Linker: Utility program which takes one or more compiled object files and combines them into an executable file or another object file.

Loader: loads the executable code into memory ,creates the program and data stack , initializes the registers and starts the code running.

```

schen@linux:~/comp6700/lec04$ strace ./hellos
execve("./hellos", ["/hellos"], 0x7fffb0c2100 /* 27 vars */) = 0
strace: [ Process PID=727737 runs in 32 bit mode. ]
arch_prctl(0x3001 /* ARCH_??? */, 0xffff3b560) = -1 EINVAL (Invalid argument)
brk(NULL) = 0x8f25000
brk(0x8f25d40) = 0x8f25d40
set_thread_area({entry_number=-1, base_addr=0x8f25840,
    limit=0x0ffff, seg_32bit=1, contents=0, read_exec_only=0,
    limit_in_pages=1, seg_not_present=0, useable=1}) = 0 (entry_number=12)
uname({sysname="Linux", nodename="linux", ...}) = 0
readlink("/proc/self/exe", "/home/schen/comp6700/lec04/hellos", 4096) = 31
brk(0x8f46d40) = 0x8f46d40
brk(0x8f47000) = 0x8f47000
access("/etc/ld.so.nohwcap", F_OK) = -1 ENOENT (No such file or directory)
mprotect(0x80e3000, 8192, PROT_READ) = 0
fstat64(1, {st_mode=S_IFCHR|0620, st_rdev=makedev(0x88, 0x1), ...}) = 0
write(1, "Hello World!\n", 13Hello World!
) = 13
exit_group(0) = ?
+++ exited with 0 +++

```

Dynamically Linked Binary

```
schen@linux:~/comp6700/lec04$ strace ./hello
execve("./hello", ["/.hello"], 0x7fff46583a90 /* 55 vars */) = 0
brk(NULL)                                = 0x560a7251c000
arch_prctl(0x3001 /* ARCH_??? */, 0x7ffd48aa4a10) = -1 EINVAL (Invalid argument)
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7f04ea3de000
access("/etc/ld.so.preload", R_OK)       = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "glibc-hwcaps/x86-64-v3/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or ...
openat(AT_FDCWD, "glibc-hwcaps/x86-64-v2/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or ...
openat(AT_FDCWD, "tls/haswell/x86_64/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or ...
openat(AT_FDCWD, "tls/haswell/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "tls/x86_64/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "tls/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "haswell/x86_64/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "haswell/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "x86_64/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "./glibc-hwcaps/x86-64-v3/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file ...
openat(AT_FDCWD, "./glibc-hwcaps/x86-64-v2/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file ...
openat(AT_FDCWD, "./tls/haswell/x86_64/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or ...
openat(AT_FDCWD, "./tls/haswell/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "./tls/x86_64/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "./tls/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "./haswell/x86_64/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or ...
openat(AT_FDCWD, "./haswell/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "./x86_64/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "./libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
newfstatat(3, "", {st_mode=S_IFREG|0644, st_size=71307, ...}, AT_EMPTY_PATH) = 0
mmap(NULL, 71307, PROT_READ, MAP_PRIVATE, 3, 0) = 0x7f04ea3cc000
close(3)                                = 0
openat(AT_FDCWD, "/lib/x86_64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\3\0>\0\1\0\0\0P\237\2\0\0\0\0"... , 832) = 832
```

Dynamically Linked Binary (continued)

```

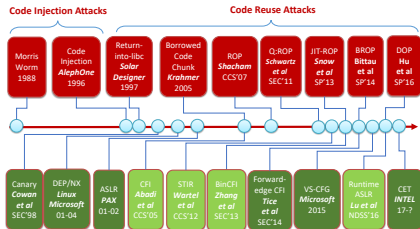
pread64(3, "\6\0\0\0\4\0\0\0@\0\0\0\0\0\0\0@\0\0\0\0\0\0\0@\0\0\0\0\0\0\0"..., 784, 64) = 784
pread64(3, "\4\0\0\0 \0\0\0\5\0\0\0GNU\0\2\0\0\300\4\0\0\0\3\0\0\0\0\0\0\0"..., 48, 848) = 48
pread64(3, "\4\0\0\0\24\0\0\0\3\0\0\0GNU\0i8\235HZ\227\223\333\350s\360\352,\223\340."..., 68, 896) = 68
newfstatat(3, "", {st_mode=S_IFREG\0644, st_size=2216304, ...}, AT_EMPTY_PATH) = 0
pread64(3, "\6\0\0\0\4\0\0\0@\0\0\0\0\0\0\0@\0\0\0\0\0\0\0@\0\0\0\0\0\0\0"..., 784, 64) = 784
mmap(NULL, 2260560, PROT_READ, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0x7f04ea000000
mmap(0x7f04ea028000, 1658880, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x28000) = ...
mmap(0x7f04ea1bd000, 360448, PROT_READ, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x1bd000) = ...
mmap(0x7f04ea215000, 24576, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x214000) = ...
mmap(0x7f04ea21b000, 52816, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = ...
close(3) = 0
mmap(NULL, 12288, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7f04ea3c9000
arch_prctl(ARCH_SET_FS, 0x7f04ea3c9740) = 0
set_tid_address(0x7f04ea3c9a10) = 9271
set_robust_list(0x7f04ea3c9a20, 24) = 0
rseq(0x7f04ea3ca0e0, 0x20, 0, 0x53053053) = 0
mprotect(0x7f04ea215000, 16384, PROT_READ) = 0
mprotect(0x560a708f6000, 4096, PROT_READ) = 0
mprotect(0x7f04ea418000, 8192, PROT_READ) = 0
prlimit64(0, RLIMIT_STACK, NULL, {rlim_cur=8192*1024, rlim_max=RLIM64_INFINITY}) = 0
munmap(0x7f04ea3cc000, 71307) = 0
newfstatat(1, "", {st_mode=S_IFCHR\0620, st_rdev=makedev(0x88, 0), ...}, AT_EMPTY_PATH) = 0
getrandom("\x1f\xee\xff\xc7\x1d\x2b\xe5\x39", 8, GRND_NONBLOCK) = 8
brk(NULL) = 0x560a7251c000
brk(0x560a7253d000) = 0x560a7253d000
write(1, "Hello world!\n", 13Hello world!
) = 13
exit_group(0) = ?
+++ exited with 0 +++

```

Dynamically Linked Binary with LD_PRELOAD

```
LD_PRELOAD=/lib/x86_64-linux-gnu/libc.so.6:/lib64/ld-linux-x86-64.so.2:linux-vdso.so.1 strace ./hello
execve("./hello", ["/./hello"], 0x7ffee8804040 /* 56 vars */) = 0
brk(NULL)                                = 0x5652a8651000
arch_prctl(0x3001 /* ARCH??? */, 0x7ffd6bcf8ac0) = -1 EINVAL (Invalid argument)
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7f433e716000
openat(AT_FDCWD, "/lib/x86_64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\0\0\0\3\0>\0\1\0\0\0P\237\2\0\0\0\0\0"... , 832) = 832
pread64(3, "\6\0\0\0\4\0\0\0@ \0\0\0\0\0\0\0\0@ \0\0\0\0\0\0\0\0@ \0\0\0\0\0\0\0\0"..., 784, 64) = 784
pread64(3, "\4\0\0\0 \0\0\0\5\0\0\0GNU\0\2\0\0\300\4\0\0\0\3\0\0\0\0\0\0\0"..., 48, 848) = 48
pread64(3, "\4\0\0\0\24\0\0\0\3\0\0\0GNU\0i8\235HZ\227\223\333\350s\360\352,\223\340."..., 68, 896) = 68
newstatat(3, "", {st_mode=S_IFREG|0644, st_size=2216304, ...}, AT_EMPTY_PATH) = 0
pread64(3, "\6\0\0\0\4\0\0\0@ \0\0\0\0\0\0\0\0@ \0\0\0\0\0\0\0\0@ \0\0\0\0\0\0\0\0"..., 784, 64) = 784
mmap(NULL, 2260560, PROT_READ, MAP_PRIVATE|MAP_DENYWRITE, 3, 0) = 0x7f433e400000
mmap(0x7f433e428000, 1658880, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x28000) = ...
mmap(0x7f433e5bd000, 360448, PROT_READ, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x1bd000) = ...
mmap(0x7f433e615000, 24576, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x214000) = ...
mmap(0x7f433e61b000, 52816, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = ...
close(3)                                = 0
access("/etc/ld.so.preload", R_OK)       = -1 ENOENT (No such file or directory)
mmap(NULL, 12288, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0x7f433e713000
arch_prctl(ARCH_SET_FS, 0x7f433e713740) = 0 set_tid_address(0x7f433e713a10)          = 9265
set_robust_list(0x7f433e713a20, 24)      = 0 rseq(0x7f433e7140e0, 0x20, 0, 0x53053053) = 0
mprotect(0x7f433e615000, 16384, PROT_READ) = 0
mprotect(0x5652a66d5000, 4096, PROT_READ) = 0
mprotect(0x7f433e750000, 8192, PROT_READ) = 0
prlimit64(0, RLIMIT_STACK, NULL, {rlim_cur=8192*1024, rlim_max=RLIM64_INFINITY}) = 0
newstatat(1, "", {st_mode=S_IFCHR|0620, st_rdev=makedev(0x88, 0), ...}, AT_EMPTY_PATH) = 0
getrandom("\x35\xaa\xb2\xf2\x9d\x9d\x62\xf5", 8, GRND_NONBLOCK) = 8
brk(NULL)                                = 0x5652a8651000
brk(0x5652a8672000)                      = 0x5652a8672000
write(1, "Hello world!\n", 13Hello world!) = 13
exit_group(0)                            = ?
+++ exited with 0 +++
```


Thank You



1



schen@auburn.edu
[schuan.github.io](https://github.com/schuan)

¹Instructor appreciates the help from Prof. Zhiqiang Lin.